

Assignment 06 Logistic Regression

```
In [1]: # Importing the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
```

```
In [2]: #loading the data
bank=pd.read_csv('bank-detail.csv')
bank
```

```
Out[2]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198
...
45206	51	technician	married	tertiary	no	825	no	no	cellular	17	nov	977
45207	71	retired	divorced	primary	no	1729	no	no	cellular	17	nov	456
45208	72	retired	married	secondary	no	5715	no	no	cellular	17	nov	1127
45209	57	blue-collar	married	secondary	no	668	no	no	telephone	17	nov	508
45210	37	entrepreneur	married	secondary	no	2971	no	no	cellular	17	nov	361

45211 rows × 17 columns

EDA

```
In [3]: bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays       45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  Y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

```
In [4]: # One-Hot Encoding of categorical variables
data1=pd.get_dummies(bank,columns=['job','marital','education','contact','poutcome','mon
data1
```

```
Out[4]:
```

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	...	month_dec	month
0	58	no	2143	yes	no	5	261	1	-1	0	...	0	
1	44	no	29	yes	no	5	151	1	-1	0	...	0	
2	33	no	2	yes	yes	5	76	1	-1	0	...	0	
3	47	no	1506	yes	no	5	92	1	-1	0	...	0	
4	33	no	1	no	no	5	198	1	-1	0	...	0	
...	
45206	51	no	825	no	no	17	977	3	-1	0	...	0	
45207	71	no	1729	no	no	17	456	2	-1	0	...	0	
45208	72	no	5715	no	no	17	1127	5	184	3	...	0	
45209	57	no	668	no	no	17	508	4	-1	0	...	0	
45210	37	no	2971	no	no	17	361	2	188	11	...	0	

45211 rows × 49 columns

```
In [5]: # To see all columns
pd.set_option("display.max.columns", None)
data1
```

Out[5]:

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	Y	job_admin.	job_
0	58	no	2143	yes	no	5	261	1	-1	0	no	0	
1	44	no	29	yes	no	5	151	1	-1	0	no	0	
2	33	no	2	yes	yes	5	76	1	-1	0	no	0	
3	47	no	1506	yes	no	5	92	1	-1	0	no	0	
4	33	no	1	no	no	5	198	1	-1	0	no	0	
...
45206	51	no	825	no	no	17	977	3	-1	0	yes	0	
45207	71	no	1729	no	no	17	456	2	-1	0	yes	0	
45208	72	no	5715	no	no	17	1127	5	184	3	yes	0	
45209	57	no	668	no	no	17	508	4	-1	0	no	0	
45210	37	no	2971	no	no	17	361	2	188	11	no	0	

45211 rows × 49 columns

In [6]:

data1.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 45211 entries, 0 to 45210
```

```
Data columns (total 49 columns):
```

#	Column	Non-Null Count	Dtype
0	age	45211 non-null	int64
1	default	45211 non-null	object
2	balance	45211 non-null	int64
3	housing	45211 non-null	object
4	loan	45211 non-null	object
5	day	45211 non-null	int64
6	duration	45211 non-null	int64
7	campaign	45211 non-null	int64
8	pdays	45211 non-null	int64
9	previous	45211 non-null	int64
10	Y	45211 non-null	object
11	job_admin.	45211 non-null	uint8
12	job_blue-collar	45211 non-null	uint8
13	job_entrepreneur	45211 non-null	uint8
14	job_housemaid	45211 non-null	uint8
15	job_management	45211 non-null	uint8
16	job_retired	45211 non-null	uint8
17	job_self-employed	45211 non-null	uint8
18	job_services	45211 non-null	uint8
19	job_student	45211 non-null	uint8
20	job_technician	45211 non-null	uint8
21	job_unemployed	45211 non-null	uint8
22	job_unknown	45211 non-null	uint8
23	marital_divorced	45211 non-null	uint8
24	marital_married	45211 non-null	uint8
25	marital_single	45211 non-null	uint8
26	education_primary	45211 non-null	uint8
27	education_secondary	45211 non-null	uint8
28	education_tertiary	45211 non-null	uint8
29	education_unknown	45211 non-null	uint8
30	contact_cellular	45211 non-null	uint8
31	contact_telephone	45211 non-null	uint8
32	contact_unknown	45211 non-null	uint8
33	poutcome_failure	45211 non-null	uint8
34	poutcome_other	45211 non-null	uint8
35	poutcome_success	45211 non-null	uint8
36	poutcome_unknown	45211 non-null	uint8
37	month_apr	45211 non-null	uint8
38	month_aug	45211 non-null	uint8
39	month_dec	45211 non-null	uint8
40	month_feb	45211 non-null	uint8
41	month_jan	45211 non-null	uint8
42	month_jul	45211 non-null	uint8
43	month_jun	45211 non-null	uint8
44	month_mar	45211 non-null	uint8
45	month_may	45211 non-null	uint8
46	month_nov	45211 non-null	uint8
47	month_oct	45211 non-null	uint8
48	month_sep	45211 non-null	uint8

```
dtypes: int64(7), object(4), uint8(38)
```

```
memory usage: 5.4+ MB
```

```
In [7]: # Custom Binary Encoding of Binary o/p variables
data1['default'] = np.where(data1['default'].astype(str).str.contains("yes"), 1, 0)
data1['housing'] = np.where(data1['housing'].astype(str).str.contains("yes"), 1, 0)
data1['loan'] = np.where(data1['loan'].astype(str).str.contains("yes"), 1, 0)
data1['Y'] = np.where(data1['Y'].astype(str).str.contains("yes"), 1, 0)
data1
```

Out[7]:

	age	default	balance	housing	loan	day	duration	campaign	pdays	previous	Y	job_admin.	job_b cc
0	58	0	2143	1	0	5	261	1	-1	0	0	0	
1	44	0	29	1	0	5	151	1	-1	0	0	0	
2	33	0	2	1	1	5	76	1	-1	0	0	0	
3	47	0	1506	1	0	5	92	1	-1	0	0	0	
4	33	0	1	0	0	5	198	1	-1	0	0	0	
...
45206	51	0	825	0	0	17	977	3	-1	0	1	0	
45207	71	0	1729	0	0	17	456	2	-1	0	1	0	
45208	72	0	5715	0	0	17	1127	5	184	3	1	0	
45209	57	0	668	0	0	17	508	4	-1	0	0	0	
45210	37	0	2971	0	0	17	361	2	188	11	0	0	

45211 rows × 49 columns

In [8]:

data1.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 45211 entries, 0 to 45210
```

```
Data columns (total 49 columns):
```

#	Column	Non-Null Count	Dtype
0	age	45211 non-null	int64
1	default	45211 non-null	int32
2	balance	45211 non-null	int64
3	housing	45211 non-null	int32
4	loan	45211 non-null	int32
5	day	45211 non-null	int64
6	duration	45211 non-null	int64
7	campaign	45211 non-null	int64
8	pdays	45211 non-null	int64
9	previous	45211 non-null	int64
10	Y	45211 non-null	int32
11	job_admin.	45211 non-null	uint8
12	job_blue-collar	45211 non-null	uint8
13	job_entrepreneur	45211 non-null	uint8
14	job_housemaid	45211 non-null	uint8
15	job_management	45211 non-null	uint8
16	job_retired	45211 non-null	uint8
17	job_self-employed	45211 non-null	uint8
18	job_services	45211 non-null	uint8
19	job_student	45211 non-null	uint8
20	job_technician	45211 non-null	uint8
21	job_unemployed	45211 non-null	uint8
22	job_unknown	45211 non-null	uint8
23	marital_divorced	45211 non-null	uint8
24	marital_married	45211 non-null	uint8
25	marital_single	45211 non-null	uint8
26	education_primary	45211 non-null	uint8
27	education_secondary	45211 non-null	uint8
28	education_tertiary	45211 non-null	uint8
29	education_unknown	45211 non-null	uint8
30	contact_cellular	45211 non-null	uint8
31	contact_telephone	45211 non-null	uint8
32	contact_unknown	45211 non-null	uint8
33	poutcome_failure	45211 non-null	uint8
34	poutcome_other	45211 non-null	uint8
35	poutcome_success	45211 non-null	uint8
36	poutcome_unknown	45211 non-null	uint8
37	month_apr	45211 non-null	uint8
38	month_aug	45211 non-null	uint8
39	month_dec	45211 non-null	uint8
40	month_feb	45211 non-null	uint8
41	month_jan	45211 non-null	uint8
42	month_jul	45211 non-null	uint8
43	month_jun	45211 non-null	uint8
44	month_mar	45211 non-null	uint8
45	month_may	45211 non-null	uint8
46	month_nov	45211 non-null	uint8
47	month_oct	45211 non-null	uint8
48	month_sep	45211 non-null	uint8

```
dtypes: int32(4), int64(7), uint8(38)
```

```
memory usage: 4.7 MB
```

Model Building

```
In [10]: # Dividing our data into input and output variables
x=pd.concat([data1.iloc[:,0:10],data1.iloc[:,11:]],axis=1)
```

```
y=x.astype('int')
y=data1.iloc[:,10]
```

```
In [11]: # Logistic regression model
classifier=LogisticRegression()
classifier.fit(x, y)
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[11]: ▼ LogisticRegression
LogisticRegression()
```

Model Predictions

```
In [12]: # Predict for x dataset
y_pred=classifier.predict(x)
y_pred
```

```
Out[12]: array([0, 0, 0, ..., 1, 0, 0])
```

```
In [13]: y_pred_df=pd.DataFrame({'actual_y':y,'y_pred_prob':y_pred})
y_pred_df
```

```
Out[13]:
```

	actual_y	y_pred_prob
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
45206	1	1
45207	1	0
45208	1	1
45209	0	0
45210	0	0

45211 rows × 2 columns

Testing Model Accuracy

```
In [14]: # Confusion Matrix for the model accuracy
confusion_matrix = confusion_matrix(y,y_pred)
```

confusion_matrix

```
Out[14]: array([[39152,   770],
        [ 4125,  1164]], dtype=int64)
```

```
In [80]: # The model accuracy is calculated by (a+d)/(a+b+c+d)
        (39156+1162)/(39156+766+4127+1162)
```

```
Out[80]: 0.8917741257658535
```


```
In [15]: # As accuracy = 0.8933, which is greater than 0.5; Thus[:,1] Threshold value>0.5=1 else
        classifier.predict_proba(x)[:,1]
```

```
Out[15]: array([0.04409006, 0.02468093, 0.01818396, ..., 0.67024598, 0.07890175,
        0.10202693])
```

```
In [16]: # ROC Curve plotting and finding AUC value
        fpr,tpr,thresholds=roc_curve(y,classifier.predict_proba(x)[:,1])
        plt.plot(fpr,tpr,color='red')
        auc=roc_auc_score(y,y_pred)

        plt.plot(fpr,tpr,color='red',label='logit model(area = %0.2f)'%auc)
        plt.plot([0,1],[0,1],'k--')
        plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
        plt.ylabel('True Positive Rate')
        plt.show()

        print('auc accuracy:',auc)
```



auc accuracy: 0.6003958996276432

```
In [ ]:
```